Adam Chen

Machine Learning Engineer | San Francisco, CA | adam.chen@email.com | linkedin.com/in/adamchen | github.com/adamchen

SUMMARY

Machine Learning Engineer with 4 years of experience designing and deploying NLP and computer vision systems at scale. Passionate about building end-to-end ML pipelines from research to production. Experienced with cloud infrastructure on AWS and GCP, containerized deployments, and large language model fine-tuning. Strong background in Python, PyTorch, and distributed training.

EXPERIENCE

Senior Machine Learning Engineer — Nuance AI, San Francisco

March 2022 – Present

- Designed and deployed a real-time document classification pipeline processing 500K+ documents per day using FastAPI and Kubernetes.
- Fine-tuned BERT and RoBERTa models on proprietary datasets, achieving a 12% improvement in entity extraction accuracy.
- Built an internal RAG system using FAISS and OpenAI embeddings to power a knowledge base assistant used by 200+ employees.
- Reduced model inference latency by 40% through ONNX export and TensorRT optimization on NVIDIA A100 GPUs.
- Led migration of ML workloads from on-premise servers to AWS SageMaker, cutting infrastructure costs by 30%.

Machine Learning Engineer — DataStream Labs, Austin TX

June 2020 – February 2022

- Built a sentiment analysis API used by 15 enterprise clients, serving 2M+ requests per month with 99.9% uptime.
- Developed a custom TF-IDF and semantic search engine for internal document retrieval over a corpus of 1M+ legal documents.
- Implemented data pipelines for text preprocessing using Apache Spark and deployed models via Docker on GCP Cloud Run.
- Collaborated with the product team to integrate NLP features into the SaaS platform using REST APIs.

EDUCATION

M.Sc. Computer Science — Stanford University, specialization in Artificial Intelligence

2018 – 2020 | GPA: 3.9/4.0

B.Sc. Mathematics & Computer Science — UC San Diego

2014 – 2018 | Graduated with honors

SKILLS

Languages: Python, SQL, Bash, C++

ML Frameworks: PyTorch, TensorFlow, Scikit-learn, HuggingFace Transformers

NLP / Search: BERT, RoBERTa, Whisper, FAISS, SentenceTransformers, TF-IDF

Cloud & DevOps: AWS (SageMaker, S3, Lambda), GCP (Cloud Run, BigQuery), Docker, Kubernetes

APIs & Backend: FastAPI, Flask, REST, gRPC

Other: Git, MLflow, Weights & Biases, Jupyter, Linux

PROJECTS

DocSearch — Open-source semantic document search engine (GitHub: 1.2K stars)

Built a document retrieval system supporting TF-IDF and dense embeddings with a Streamlit UI and FastAPI backend.

MedNER — Named entity recognition for clinical notes
Fine-tuned BioBERT on a custom annotated dataset of 10K clinical notes, achieving 91% F1 on medical entity extraction.